

```
In [2]: %%javascript
/*****
Known Mathjax Issue with Chrome - a rounding issue adds a border
https://github.com/mathjax/MathJax/issues/1300
A quick hack to fix this based on stackoverflow discussions:
http://stackoverflow.com/questions/34277967/chrome-rendering-mat
*****/
$($('mathspan').css("border-left-color", "transparent"))
<IPython.core.display.Javascript object>
```

```
In [3]: %reload_ext autoreload
%autoreload 2
```

MIDS - w261 Machine Learning At Scale

Course Lead: Dr James G. Shanahan (**email** Jimi via James.Shanahan AT gmail.com)

Assignment - HW3

Name: Chris Caudill

Class: MIDS w261 Spring 2017 Group 1

Email: CSCAUDILL@iSchool.Berkeley.edu

StudentId 3032134574 **End of StudentId**

Week: 3

NOTE: please replace 1234567 with your student id above

Due Time: HW is due the Tuesday of the following week by 8AM (West coast time).
I.e., Tuesday, Jan 31, 2017 in the case of this homework.

Table of Contents

1. [HW Intructions](#)
2. [HW Problems](#)
 - 3.1. [HW3.1](#)
 - 3.2. [HW3.2](#)
 - 3.3 [HW3.3](#)
 - 3.4 [HW3.4](#)
 - 3.5 [HW3.5](#)
 - 3.6 [HW3.6](#)
 - 3.7 [HW3.7](#)
 - 3.8 [HW3.8](#)

1 Instructions

[Back to Table of Contents](#)

MIDS UC Berkeley, Machine Learning at Scale DATSCIW261 ASSIGNMENT #3

Version 2017-26-1

IMPORTANT

This homework can be completed locally on your computer

=== INSTRUCTIONS for SUBMISSIONS ===

Follow the instructions for submissions carefully.

NEW: Going forward, each student will have a HW-<user> repository for all assignments.

Click this link to enable you to create a github repo within the MIDS261 Classroom:

<https://classroom.github.com/assignment-invitations>

[/3b1d6c8e58351209f9dd865537111ff8 \(https://classroom.github.com/assignment-invitations/3b1d6c8e58351209f9dd865537111ff8\)](https://classroom.github.com/assignment-invitations/3b1d6c8e58351209f9dd865537111ff8)

and follow the instructions to create a HW repo.

Push the following to your HW github repo into the master branch:

- Your local HW3 directory. Your repo file structure should look like this:

```
HW-<user>
  --HW3
    |__MIDS-W261-HW-03-<Student_id>.ipnb
    |__MIDS-W261-HW-03-<Student_id>.pdf
    |__some other hw3 file
  --HW4
    |__MIDS-W261-HW-04-<Student_id>.ipnb
    |__MIDS-W261-HW-04-<Student_id>.pdf
    |__some other hw4 file
  etc..
```

HW3.0.

1. How do you merge two sorted lists/arrays of records of the form [key, value]?
2. Where is this used in Hadoop MapReduce? [Hint within the shuffle]
3. What is a combiner function in the context of Hadoop?
4. Give an example where it can be used and justify why it should be used in the context of this problem.
5. What is the Hadoop shuffle?

HW3.0 Answers

1. With the form [key, value], we merge two sorted lists using Priority Queue.
This is initially done by establishing 3 pointers: One at the start of each sorted list, and one at start of the merged list. We start by picking off the smallest element of the two sorted lists and add them to the merged list, incrementing the pointers along the way until we reach the end of the two sorted lists.
2. sorting of merged lists occurs in the combiners
3. A combiner is used to merge the intermediary lists generated from the mappers.
4. In the context of our wordcount example, the mappers will intake records and establish an intermediary counts for each word that occurs. Although the mappers can technically perform local aggregation using arrays, a single word could still occur across multiple records or mappers. The combiner is used in this context to take those intermediary arrays and merge them into a single list. This is useful because we will be utilizing less memory by maintaining fewer lists of terms/counts.
5. The Hadoop shuffle is the entire process that happens between the mapper output and the reducer input. This includes:
 - a. Partition, sort, combine
 - b. Mergesort
 - c. Send to reducer
 - d. Mergesort partition files
 - e. Stream to reducer

Type *Markdown* and LaTeX: α^2

HW3.1 consumer complaints dataset: Use Counters to do EDA (exploratory data analysis and to monitor progress)

Counters are lightweight objects in Hadoop that allow you to keep track of system progress in both the map and reduce stages of processing. By default, Hadoop defines a number of standard counters in "groups"; these show up in the jobtracker webapp, giving you information such as "Map input records", "Map output records", etc.

While processing information/data using MapReduce job, it is a challenge to monitor the progress of parallel threads running across nodes of distributed clusters. Moreover, it is also complicated to distinguish between the data that has been processed and the data which is yet to be processed. The MapReduce Framework offers a provision of user-defined Counters, which can be effectively utilized to monitor the progress of data across nodes of distributed clusters.

Use the Consumer Complaints Dataset provide here to complete this question:

https://www.dropbox.com/s/vbalm3yva2rr86m/Consumer_Complaints.csv?dl=0

The consumer complaints dataset consists of diverse consumer complaints, which have been reported across the United States regarding various types of loans. The dataset consists of records of the form:

Complaint ID,Product,Sub-product,Issue,Sub-issue,State,ZIP code,Submitted via,Date received,Date sent to company,Company,Company response,Timely response?,Consumer disputed?

Here's is the first few lines of the of the Consumer Complaints Dataset:

```
Complaint ID,Product,Sub-product,Issue,Sub-issue,State,ZIP
code,Submitted via,Date received,Date sent to company,C
ompany,Company response,Timely response?,Consumer dispute
d?
1114245,Debt collection,Medical,Disclosure verification o
f debt,Not given enough info to verify debt,FL,32219,Web,
11/13/2014,11/13/2014,"Choice Recovery, Inc.",Closed with
explanation,Yes,
1114488,Debt collection,Medical,Disclosure verification o
f debt,Right to dispute notice not received,TX,75006,Web,
11/13/2014,11/13/2014,"Expert Global Solutions, Inc.",In
progress,Yes,
1114255,Bank account or service,Checking account,Deposits
and withdrawals,NY,11102,Web,11/13/2014,11/13/2014,"FNT
```

```
In [2]: # Create HDFS directories
mkdir ConsumerComplaints
mkdir: cannot create directory `ConsumerComplaints': File exists
```

```
In [20]: # Put the data into HDFS
!curl 'https://www.dropbox.com/s/vbalm3yva2rr86m/Consumer_Complaints/ConsumerComplaints/'
% Total      % Received % Xferd  Average Speed   Time    Time
      Time Current              Dload  Upload   Total   Spent
      Left  Speed
  0     0    0     0    0    0     0      0  --:--:--  --:--:--
--:--:--    0
100 48.5M  100 48.5M    0     0 4774k      0  0:00:10  0:00:10
--:--:-- 5558k
```

```
In [47]: %%writefile complaintCountsMapper.py
#!/usr/bin/env python

# # START STUDENT CODE HW31MAPPER

import sys
from collections import defaultdict

# Set up counters to monitor/understand the number of times a map

prods = defaultdict(int)

for line in sys.stdin:
    if 'debt collection' in line.lower():
        prods['debt collection'] += 1
        sys.stderr.write("reporter:counter:HW3.1 Mapper Debt Col

    elif 'mortgage' in line.lower():
        prods['mortgage'] += 1
        sys.stderr.write("reporter:counter:HW3.1 Mapper Mortgage

    else:
        prods['other'] += 1
        sys.stderr.write("reporter:counter:HW3.1 Mapper Other Co

for key in sorted(prods):
    print '%s\t%s' % (key, prods[key])

# END STUDENT CODE HW31MAPPER

Overwriting complaintCountsMapper.py
```

Not required to create reducer, but for the sake of continuity, I've created one

```
In [40]: %%writefile complaintCountsReducer.py
#!/usr/bin/env python

# START STUDENT CODE HW31REDUCER

import sys
# Set up counters to monitor/understand the number of times a re
sys.stderr.write("reporter:counter:HW3.1 Reducer Counters,Calls,1

d = {}

# For each line from the mapper output, split the three attributes
# Spam/Ham classification, key term, value count
for line in sys.stdin:
    key,value = line.split("\t")

    # check if docClass + key term combination exist in dictionary
    # if so, update the term count value
    if key in d:
        d[key] += int(value)

    # if combination does not exist, add it to the dictionary
    else:
        d[key] = int(value)

# for each value in the dictionary, print the 3 attributes in a
# way that they can later be sorted appropriately
for key, value in d.iteritems():
    print '%s\t%s' % (key, value)

# END STUDENT CODE HW31REDUCER
```

Overwriting complaintCountsReducer.py

Update Security of Mapper and Reducer

```
In [41]: !chmod a+x complaintCountsMapper.py
!chmod a+x complaintCountsReducer.py
```

Execute Hadoop calls

```
In [46]: # Hadoop command
# START STUDENT CODE HW31HADOOP

!hdfs dfs -rm ccd.csv
!hdfs dfs -copyFromLocal ccd.csv
!hdfs dfs -rm -r cc-output

!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.6.0-cdh5.8.0.jar \
  -files complaintCountsMapper.py,complaintCountsReducer.py \
  -mapper complaintCountsMapper.py \
  -reducer complaintCountsReducer.py \
  -combiner complaintCountsReducer.py \
  -input ccd.csv \
  -output cc-output \
  -numReduceTasks 1

# END STUDENT CODE HW31HADOOP

Deleted ccd.csv
Deleted cc-output
packageJobJar: [] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.8.0.jar] /tmp/streamjob5158676712758093896.jar tmpDir=null
17/01/28 13:33:58 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032
17/01/28 13:33:58 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032
17/01/28 13:33:58 INFO mapred.FileInputFormat: Total input paths to process : 1
17/01/28 13:33:58 INFO mapreduce.JobSubmitter: number of splits :2
17/01/28 13:33:58 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1484544278544_0363
17/01/28 13:33:59 INFO impl.YarnClientImpl: Submitted application application_1484544278544_0363
17/01/28 13:33:59 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1484544278544_0363/ (http://quickstart.cloudera:8088/proxy/application_1484544278544_0363/)
```

Check output counts

```
In [48]: !hdfs dfs -cat cc-output/part-0000* > cc_counts.txt
!head cc_counts.txt
debt collection 44372
other 142745
mortgage 125796
```


HW 3.2 Analyze the performance of your Mappers, Combiners and Reducers using Counters

For this brief study the Input file will be one record (the next line only):

```
foo foo quux labs foo bar quux
```

3.2.A

Perform a word count analysis of this single record dataset using a Mapper and Reducer based WordCount (i.e., no combiners are used here) using user defined Counters to count up how many times the mapper and reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing this word count job. The answer should be 1 and 4 respectively. Please explain.

3.2.B

Please use multiple mappers and reducers for these jobs (at least 2 mappers and 2 reducers). Perform a word count analysis of the Issue column of the Consumer Complaints Dataset using a Mapper and Reducer based WordCount (i.e., no combiners used anywhere) using user defined Counters to count up how many time the mapper and reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing your word count job.

3.2.C

Perform a word count analysis of the Issue column of the Consumer Complaints Dataset using a Mapper, Reducer, and standalone combiner (i.e., not an in-memory combiner) based WordCount using user defined Counters to count up how many time the mapper, combiner, reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing your word count job.

Using a single reducer:

- What are the top 50 most frequent terms in your word count analysis?
- Present the top 50 terms and their frequency and their relative frequency. If there are ties please sort the tokens in alphanumeric/string order.
- Present bottom 10 tokens (least frequent items).

NOTE: You can use: `WORD_RE = re.compile(r"[\w']+")` to tokenize.

3.2.A SOLUTION

```
In [70]: %%writefile mapper3.2.A.py
#!/usr/bin/env python
# START STUDENT CODE HW32AMAPPER

import sys
from collections import defaultdict

# Set up counters to monitor/understand the number of times a map
sys.stderr.write("reporter:counter:HW3.2.A Mapper Counters,Calls

wordCounts = defaultdict(int)

for line in sys.stdin:
    words = line.split()
    for word in words:
        print '%s\t%s' % (word, 1)
#         if word in wordCounts:
#             wordCounts[word] += 1
#         else:
#             wordCounts[word] += 1

# for key in sorted(wordCounts):
#     print '%s\t%s' % (key, wordCounts[key])

# END STUDENT CODE HW32AMAPPER
```

Overwriting mapper3.2.A.py

```
In [71]: %%writefile reducer3.2.A.py
#!/usr/bin/env python
# START STUDENT CODE HW32AREducer

import sys
# Set up counters to monitor/understand the number of times a re
sys.stderr.write("reporter:counter:HW3.2.A Reducer Counters,Call:

d = {}

# For each line from the mapper output, split the three attribute
# Spam/Ham classification, key term, value count
for line in sys.stdin:
    key,value = line.split("\t")

    # check if docClass + key term combination exist in dictionary
    # if so, update the term count value
    if key in d:
        d[key] += int(value)

    # if combination does not exist, add it to the dictionary
    else:
        d[key] = int(value)

# for each value in the dictionary, print the 3 attributes in a
# way that they can later be sorted appropriately
for key, value in d.iteritems():
    print '%s\t%s' % (key, value)

# END STUDENT CODE HW32AREducer
```

Overwriting reducer3.2.A.py

Create input file

```
In [61]: echo "foo foo quux labz foo bar quux" > /HW3.2.txt
```

```
In [82]: # Hadoop command
# START STUDENT CODE HW32AHAD00P

#update security on mapper/reducer
!chmod a+x mapper3.2.A.py
!chmod a+x reducer3.2.A.py

!hdfs dfs -rm HW3.2.txt
!hdfs dfs -copyFromLocal HW3.2.txt
!hdfs dfs -rm -r HW3.2.A-output

!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2
    -files mapper3.2.A.py,reducer3.2.A.py \
    -mapper mapper3.2.A.py \
    -reducer reducer3.2.A.py \
    -input HW3.2.txt \
    -output HW3.2.A-output \
    -numReduceTasks 4

# END STUDENT CODE HW32AHAD00P
```

Deleted HW3.2.txt

rm: `HW3.2.A-output': No such file or directory

packageJobJar: [] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.8.0.jar] /tmp/streamjob4205658107599752123.jar tmpDir=null

17/01/28 14:36:41 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032

17/01/28 14:36:41 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032

17/01/28 14:36:42 INFO mapred.FileInputFormat: Total input paths to process : 1

17/01/28 14:36:42 INFO mapreduce.JobSubmitter: number of splits : 2

17/01/28 14:36:42 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1484544278544_0371

17/01/28 14:36:42 INFO impl.YarnClientImpl: Submitted application application_1484544278544_0371

17/01/28 14:36:42 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1484544278544_0371/ (http://quickstart.cloudera:8088/proxy/application_1484544278544_0371/)

```
In [107]: # INSERT SCREENSHOT OF JOB TRACKER UI COUNTERS
```

```
from IPython.display import Image, HTML
Image('HW32.png')
```

Out[107]:

HW3.2.A Mapper Counters	Name	Map	Reduce	Total
Calls		2	0	2
HW3.2.A Reducer Counters	Name	Map	Reduce	Total
Calls		0	4	4

3.2.A EXPLANATION

For mappers, I accepted the defaults and used 2. For reducers, I specified for 4 to be used. These numbers will change based on the settings and configuration that is used within the Hadoop Streaming calls.

3.2.B SOLUTION

3.2.B Mapper

```
In [108]: %%writefile mapper3.2.B.py
          #!/usr/bin/env python
          # START STUDENT CODE HW32BMAPPER

          import sys
          import csv
          from collections import defaultdict
          import re

          # Set up counters to monitor/understand the number of times a map
          sys.stderr.write("reporter:counter:HW3.2.B Mapper Counters,Calls

          issueList = defaultdict(int)

          ccd = csv.reader(sys.stdin)
          for row in ccd:
              issue = row[3]
              if len(issue) > 0 and issue != 'Issue':
                  words = re.findall(r'[a-z]+', issue.lower())
                  for word in words:
                      if word in issueList:
                          issueList[word] +=1
                      else:
                          issueList[word] = 1

          for key in sorted(issueList):
              print '%s\t%s' % (key, issueList[key])

          # END STUDENT CODE HW32BMAPPER
```

Overwriting mapper3.2.B.py

3.2.B Reducer

```
In [109]: %%writefile reducer3.2.B.py
#!/usr/bin/env python
# START STUDENT CODE HW32BREDUCER

import sys
# Set up counters to monitor/understand the number of times a re
sys.stderr.write("reporter:counter:HW3.2.B Reducer Counters,Call:

d = {}

# For each line from the mapper output, split the three attribute
# Spam/Ham classification, key term, value count
for line in sys.stdin:
    key,value = line.split("\t")

    # check if docClass + key term combination exist in dictionary
    # if so, update the term count value
    if key in d:
        d[key] += int(value)

    # if combination does not exist, add it to the dictionary
    else:
        d[key] = int(value)

# for each value in the dictionary, print the 3 attributes in a
# way that they can later be sorted appropriately
for key, value in d.iteritems():
    print '%s\t%s' % (key, value)

# END STUDENT CODE HW32BREDUCER
```

Overwriting reducer3.2.B.py

3.2.B Hadoop Calls

```

In [110]: # Hadoop command
# START STUDENT CODE HW32BHAD00P

!chmod a+x mapper3.2.B.py
!chmod a+x reducer3.2.B.py

!hdfs dfs -rm ccd.csv
!hdfs dfs -copyFromLocal ccd.csv
!hdfs dfs -rm -r HW3.2.B-output

!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.6.0-cdh5.8.0.jar \
  -files mapper3.2.B.py,reducer3.2.B.py \
  -reducer reducer3.2.B.py \
  -mapper mapper3.2.B.py \
  -input ccd.csv \
  -output HW3.2.B-output \
  -numReduceTasks 4

# END STUDENT CODE HW32BHAD00P
Deleted ccd.csv
Deleted HW3.2.B-output
packageJobJar: [] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.8.0.jar] /tmp/streamjob3908761271563178578.jar tmpDir=null
17/01/29 15:30:05 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
17/01/29 15:30:06 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
17/01/29 15:30:06 INFO mapred.FileInputFormat: Total input paths to process : 1
17/01/29 15:30:06 INFO mapreduce.JobSubmitter: number of splits :2
17/01/29 15:30:06 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1484544278544_0417
17/01/29 15:30:07 INFO impl.YarnClientImpl: Submitted application application_1484544278544_0417
17/01/29 15:30:07 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1484544278544_0417/ (http://quickstart.cloudera:8088/proxy/application_1484544278544_0417/)

```

3.2.B Answer

```
In [111]: # 3.2.B OUTPUT/ANSWER
!hdfs dfs -cat HW3.2.B-output/part-0000* > HW3.2.B_counts.txt
!cat HW3.2.B_counts.txt
owed      17972
caused    5663
sale       139
acct       163
your      3844
payoff    1155
incorrect      29133
management    16205
delinquent    1061
for           929
transfer      597
the           6248
to            8401
unable       8178
issue        1098
checks       75
available    274
applied      139
workout      350
use          274
```

```
In [129]: # INSERT SCREENSHOT OF JOB TRACKER UI COUNTERS
from IPython.display import Image, HTML
Image('HW32B.png')
```

Out[129]:

HW3.2.B Mapper Counters	Name	Map	Reduce	Total
Calls		2	0	2
HW3.2.B Reducer Counters	Name	Map	Reduce	Total
Calls		0	4	4

3.2.C SOLUTION

3.2.C Mapper


```
In [117]: %%writefile mapper3.2.C.py
#!/usr/bin/env python
# START STUDENT CODE HW32CMAPPER

import sys
import csv
import re

# Set up counters to monitor/understand the number of times a map
sys.stderr.write("reporter:counter:HW3.2.B Mapper Counters,Calls

ccd = csv.reader(sys.stdin)
for row in ccd:
    issue = row[3]
    if len(issue) > 0 and issue != 'Issue':
        words = re.findall(r'[a-z]+', issue.lower())
        for word in words:
            print '%s\t%s' % (word, 1)

# END STUDENT CODE HW32CMAPPER
Overwriting mapper3.2.C.py
```

3.2.C Combiner

```
In [122]: %%writefile combiner3.2.C.py
#!/usr/bin/env python
# START STUDENT CODE HW32CCOMBINER

import sys
from collections import defaultdict

# Set up counters to monitor/understand the number of times a red
sys.stderr.write("reporter:counter:HW3.2.C Combiner Counters,Calls

issueDict = defaultdict(int)

for line in sys.stdin:
    word,count = line.split("\t")
    if word in issueDict:
        issueDict[word] += int(count)
    else:
        issueDict[word] = int(count)

for key in sorted(issueDict):
    print '%s\t%s' % (key, issueDict[key])

# END STUDENT CODE HW32CCOMBINER
Overwriting combiner3.2.C.py
```

3.2.C Reducer

```
In [123]: %%writefile reducer3.2.C.py
          #!/usr/bin/env python
          # START STUDENT CODE HW32CREDCER

          import sys
          # Set up counters to monitor/understand the number of times a red
          sys.stderr.write("reporter:counter:HW3.2.C Reducer Counters,Call

          d = {}

          # For each line from the mapper output, split the three attributes
          # Spam/Ham classification, key term, value count
          for line in sys.stdin:
              key,value = line.split("\t",1)

              # check if docClass + key term combination exist in dictionary
              # if so, update the term count value
              if key in d:
                  d[key] += int(value)

              # if combination does not exist, add it to the dictionary
              else:
                  d[key] = int(value)

          # for each value in the dictionary, print the 3 attributes in a
          # way that they can later be sorted appropriately
          for key, value in d.iteritems():
              print '%s\t%s' % (key, value)

          # END STUDENT CODE HW32CREDCER
Overwriting reducer3.2.C.py
```

3.2.C Hadoop Calls

```
In [124]: # Hadoop command
# START STUDENT CODE HW32CHAD00P

!chmod a+x mapper3.2.C.py
!chmod a+x combiner3.2.C.py
!chmod a+x reducer3.2.C.py

!hdfs dfs -rm ccd.csv
!hdfs dfs -copyFromLocal ccd.csv
!hdfs dfs -rm -r HW3.2.C-output

!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.6.0-cdh5.8.0.jar \
  -files mapper3.2.C.py,combiner3.2.C.py,reducer3.2.C.py \
  -reducer reducer3.2.C.py \
  -combiner combiner3.2.C.py \
  -mapper mapper3.2.C.py \
  -input ccd.csv \
  -output HW3.2.C-output \
  -numReduceTasks 1

# END STUDENT CODE HW32CHAD00P
```

```
Deleted ccd.csv
Deleted HW3.2.C-output
packageJobJar: [] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.8.0.jar] /tmp/streamjob6631041556751413680.jar tmpDir=null
17/01/29 15:54:13 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
17/01/29 15:54:13 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
17/01/29 15:54:14 INFO mapred.FileInputFormat: Total input paths to process : 1
17/01/29 15:54:14 INFO mapreduce.JobSubmitter: number of splits : 2
17/01/29 15:54:14 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1484544278544_0420
17/01/29 15:54:14 INFO impl.YarnClientImpl: Submitted application application_1484544278544_0420
17/01/29 15:54:14 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1484544278544_0420/ (http://quickstart.cloudera:8088/proxy/application_1484544278544_0420/)
```

3.2.C Answer

```
In [125]: # 3.2.C OUTPUT/ANSWER
!hdfs dfs -cat HW3.2.C-output/part-0000* > HW3.2.C_counts.txt
!cat HW3.2.C_counts.txt
```

```
unsolicited      640
foreclosure      70487
being    5663
caused    5663
scam      566
embezzlement    3276
report    34903
attempts    17972
settlement    4350
underwriting    2774
issues    538
payoff    1155
delinquent    1061
credited      92
deposits    10555
amt          71
late        1797
to          8401
other       7886
money       2620
```

```
In [128]: # INSERT SCREENSHOT OF JOB TRACKER UI COUNTERS
from IPython.display import Image, HTML
Image('HW32C.png')
```

```
Out[128]:
```

	virtual memory (bytes)	snapshot	3123123120	1370918400	4096045520	
HW3.2.B Mapper Counters	Name	^	Map	◊	Reduce	◊
	Calls		2	0	2	Total
HW3.2.C Combiner Counters	Name	^	Map	◊	Reduce	◊
	Calls		2	0	2	Total
HW3.2.C Reducer Counters	Name	^	Map	◊	Reduce	◊
	Calls		0	1	1	Total

3.2.1

Using **2 reducers**: What are the top **50 most frequent terms** in your word count analysis?

Present the top 50 terms and their frequency and their relative frequency. Present the top 50 terms and their frequency and their relative frequency. If there are ties please sort the tokens in alphanumeric/string order. Present bottom 10 tokens (least frequent items). Please **use a combiner**.

START STUDENT CODE HW321 (INSERT CELLS BELOW AS NEEDED)

3.2.C Frequency Mapper

```
In [152]: %%writefile frequencies_mapper3.2.C.py
          #!/usr/bin/env python
          # START STUDENT CODE HW32CFREQMAPPER

          import sys, csv, re

          # Set up counters to monitor/understand the number of times a map
          sys.stderr.write("reporter:counter:HW3.2.C Mapper Counters,Calls

          ccd = csv.reader(sys.stdin)
          for row in ccd:
              issue = row[3]
              if len(issue) > 0 and issue != 'Issue':
                  words = re.findall(r'[a-z]+', issue.lower())
                  for word in words:
                      print '%s\t%s\t%s' % (word, 1,10)

          # END STUDENT CODE HW32CFREQMAPPER
```

Overwriting frequencies_mapper3.2.C.py

3.2.C Frequency Reducer

```
In [165]: %%writefile frequencies_reducer3.2.C.py
#!/usr/bin/env python
# START STUDENT CODE HW32CFREQREDUCER

import sys
# Set up counters to monitor/understand the number of times a re
sys.stderr.write("reporter:counter:HW3.2.C Reducer Counters,Call:

d = {}
total = 0
# For each line from the mapper output, split the three attribute
# Spam/Ham classification, key term, value count
for line in sys.stdin:
    key,value,x = line.split("\t",2)

    # check if docClass + key term combination exist in dictionary
    # if so, update the term count value
    if key in d:
        d[key] += int(value)

    # if combination does not exist, add it to the dictionary
    else:
        d[key] = int(value)

for key, value in d.iteritems():
    total += int(value)

# for each value in the dictionary, print the 3 attributes in a
# way that they can later be sorted appropriately
for key, value in d.iteritems():
    freq = 0
    freq = round((float(value) / float(total)),4)
    print '%s\t%s\t%s' % (key, value, freq)

# END STUDENT CODE HW32CFREQREDUCER
```

Overwriting frequencies_reducer3.2.C.py

3.2.C Hadoop Calls

```
In [266]: # Hadoop command
# START STUDENT CODE HW32CFREQHAD00P

!chmod a+x frequencies_mapper3.2.C.py
!chmod a+x frequencies_reducer3.2.C.py

!hdfs dfs -rm ccd.csv
!hdfs dfs -copyFromLocal ccd.csv
!hdfs dfs -rm -r HW3.2.C.freq-output

!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.6.0-cdh5.8.0.jar \
  -files frequencies_mapper3.2.C.py,frequencies_reducer3.2.C.py \
  -reducer frequencies_reducer3.2.C.py \
  -combiner frequencies_reducer3.2.C.py \
  -mapper frequencies_mapper3.2.C.py \
  -input ccd.csv \
  -output HW3.2.C.freq-output \
  -numReduceTasks 2

# END STUDENT CODE HW32CFREQHAD00P

Deleted ccd.csv
Deleted HW3.2.C.freq-output
packageJobJar: [] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.8.0.jar] /tmp/streamjob4385493534662397680.jar tmpDir=null
17/01/30 21:21:52 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
17/01/30 21:21:52 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
17/01/30 21:21:53 INFO mapred.FileInputFormat: Total input paths to process : 1
17/01/30 21:21:53 INFO mapreduce.JobSubmitter: number of splits :2
17/01/30 21:21:53 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1484544278544_0457
17/01/30 21:21:53 INFO impl.YarnClientImpl: Submitted application application_1484544278544_0457
17/01/30 21:21:53 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1484544278544_0457/ (http://quickstart.cloudera:8088/proxy/application_1484544278544_0457/)
```

3.2.C Answer

50 Most Frequent Terms

```
In [267]: # 3.2.C OUTPUT/ANSWER
!hdfs dfs -cat HW3.2.C.freq-output/part-0000* > HW3.2.C.freq_cou
!cat HW3.2.C.freq_counts.txt | sort -k2,2nr | head -n50
```

loan	119630	0.1707
collection	72394	0.1033
foreclosure	70487	0.1047
modification	70487	0.1006
account	57448	0.0853
credit	55251	0.0821
or	40508	0.0602
payments	39993	0.0594
escrow	36767	0.0546
servicing	36767	0.0525
report	34903	0.0498
incorrect	29133	0.0433
information	29069	0.0415
on	29069	0.0432
debt	27874	0.0414
closing	19000	0.0282
not	18477	0.0274
attempts	17972	0.0256
collect	17972	0.0256
cost	17072	0.0256

10 Least Frequent Terms

```
In [173]: # 3.2.C OUTPUT/ANSWER
!hdfs dfs -cat HW3.2.C.freq-output/part-0000* > HW3.2.C.freq_cou
!cat HW3.2.C.freq_counts.txt | sort -k2,2nr | head -n10
```

disclosures	64	0.0001
missing	64	0.0001
amt	71	0.0001
day	71	0.0001
checks	75	0.0001
convenience	75	0.0001
credited	92	0.0001
payment	92	0.0001
amount	98	0.0001
apply	118	0.0002

END STUDENT CODE HW321

HW3.3. Shopping Cart Analysis

Product Recommendations: The action or practice of selling additional products or services to existing customers is called cross-selling. Giving product recommendation is one of the examples of cross-selling that are frequently used by online retailers. One simple method to give product recommendations is to recommend products that are frequently browsed together by the customers.

For this homework use the online browsing behavior dataset located at:

`https://www.dropbox.com/s/zlfyiwa70poqg74/ProductPurchaseData.txt?dl=0`

Each line in this dataset represents a browsing session of a customer. On each line, each string of 8 characters represents the id of an item browsed during that session. The items are separated by spaces.

Here are the first few lines of the ProductPurchaseData FRO11987 ELE17451
ELE89019 SNA90258 GRO99222 GRO99222 GRO12298 FRO12685 ELE91550
SNA11465 ELE26917 ELE52966 FRO90334 SNA30755 ELE17451 FRO84225
SNA80192 ELE17451 GRO73461 DAI22896 SNA99873 FRO86643 ELE17451
ELE37798 FRO86643 GRO56989 ELE23393 SNA11465 ELE17451 SNA69641
FRO86643 FRO78087 SNA11465 GRO39357 ELE28573 ELE11375 DAI54444

Do some exploratory data analysis of this dataset guided by the following questions:.

How many unique items are available from this supplier?

Using a single reducer: Report your findings such as number of unique products; largest basket; report the top 50 most frequently purchased items, their frequency, and their relative frequency (break ties by sorting the products alphabetical order) etc. using Hadoop Map-Reduce.

START STUDENT CODE HW33 (INSERT CELLS BELOW AS NEEDED)

In [174]: `!curl -H"Host: www.dropbox.com" -s -L -u "fujus70@gmail.com" https://www.dropbox.com/Products/PurchaseData.txt`

```

% Total    % Received % Xferd  Average Speed   Time    Time
      Time Current                      Dload  Upload   Total   Spent
      Left  Speed
  0     0    0     0    0    0     0      0  --:--:--  --:--:--
--:--:--    0
100 3377k  100 3377k    0    0 1371k      0  0:00:02  0:00:02
--:--:-- 2356k

```

In [175]: `!head ProductPurchaseData.txt`

```

FR011987 ELE17451 ELE89019 SNA90258 GR099222
GR099222 GR012298 FR012685 ELE91550 SNA11465 ELE26917 ELE52966
FR090334 SNA30755 ELE17451 FR084225 SNA80192
ELE17451 GR073461 DAI22896 SNA99873 FR086643
ELE17451 ELE37798 FR086643 GR056989 ELE23393 SNA11465
ELE17451 SNA69641 FR086643 FR078087 SNA11465 GR039357 ELE28573
ELE11375 DAI54444
ELE17451 GR073461 DAI22896 SNA99873 FR018919 DAI50921 SNA80192
GR075578
ELE17451 ELE59935 FR018919 ELE23393 SNA80192 SNA85662 SNA91554
DAI22177
ELE17451 SNA69641 FR018919 SNA90258 ELE28573 ELE11375 DAI14125
FR078087
ELE17451 GR073461 DAI22896 SNA80192 SNA85662 SNA90258 DAI46755
FR081176 ELE66810 DAI49199 DAI91535 GR094758 ELE94711 DAI22177
ELE17451 SNA69641 DAI91535 GR094758 GR099222 FR076833 FR081176
SNA80192 DAI54690 ELE37798 GR056989

```

3.3 Mapper

```

In [176]: %%writefile mapper3.3.py
          #!/usr/bin/env python

          import sys

          # Set up counters to monitor/understand the number of times a map
          sys.stderr.write("reporter:counter:HW3.3 Mapper Counters,Calls,1")

          for line in sys.stdin:
              prods = line.split()
              for item in prods:
                  print '%s\t%s\t%s' % (item, 1, 1)

```

Writing mapper3.3.py

3.3 Reducer

```
In [179]: %%writefile reducer3.3.py
#!/usr/bin/env python

import sys
# Set up counters to monitor/understand the number of times a re
sys.stderr.write("reporter:counter:HW3.3 Reducer Counters,Calls,")

d = {}
total = 0
# For each line from the mapper output, split the three attributes
# Spam/Ham classification, key term, value count
for line in sys.stdin:
    key,value,x = line.split("\t",2)

    # check if docClass + key term combination exist in dictionary
    # if so, update the term count value
    if key in d:
        d[key] += int(value)

    # if combination does not exist, add it to the dictionary
    else:
        d[key] = int(value)

for key, value in d.iteritems():
    total += int(value)

# for each value in the dictionary, print the 3 attributes in a
# way that they can later be sorted appropriately
for key, value in d.iteritems():
    freq = 0
    freq = round((float(value) / float(total)),4)
    print '%s\t%s\t%s' % (key, value, freq)
```

Overwriting reducer3.3.py

3.3 Hadoop Calls

In [185]: *# Hadoop command*

```
!chmod a+x mapper3.3.py
!chmod a+x reducer3.3.py

!hdfs dfs -rm ProductPurchaseData.txt
!hdfs dfs -copyFromLocal ProductPurchaseData.txt
!hdfs dfs -rm -r HW3.3-output

!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.6.0-cdh5.8.0.jar \
  -files mapper3.3.py,reducer3.3.py \
  -reducer reducer3.3.py \
  -combiner reducer3.3.py \
  -mapper mapper3.3.py \
  -input ProductPurchaseData.txt \
  -output HW3.3-output \
  numReduceTasks 1
```

Deleted ProductPurchaseData.txt

Deleted HW3.3-output

packageJobJar: [] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.8.0.jar] /tmp/streamjob6402019114965539630.jar tmpDir=null

17/01/29 17:08:42 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032

17/01/29 17:08:43 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032

17/01/29 17:08:43 INFO mapred.FileInputFormat: Total input paths to process : 1

17/01/29 17:08:43 INFO mapreduce.JobSubmitter: number of splits :2

17/01/29 17:08:43 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1484544278544_0432

17/01/29 17:08:44 INFO impl.YarnClientImpl: Submitted application application_1484544278544_0432

17/01/29 17:08:44 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1484544278544_0432/ (http://quickstart.cloudera:8088/proxy/application_1484544278544_0432/)

3.3 ANSWER

Top 50 Items

```
In [187]: # 3.3 OUTPUT/ANSWER
!hdfs dfs -cat HW3.3-output/part-0000* > HW3.3counts.txt
!cat HW3.3counts.txt | sort -k2 -nr | head -n50
```

DAI62779	6667	0.0175
FR040251	3881	0.0102
ELE17451	3875	0.0102
GR073461	3602	0.0095
SNA80324	3044	0.008
ELE32164	2851	0.0075
DAI75645	2736	0.0072
SNA45677	2455	0.0064
FR031317	2330	0.0061
DAI85309	2293	0.006
ELE26917	2292	0.006
FR080039	2233	0.0059
GR021487	2115	0.0056
SNA99873	2083	0.0055
GR059710	2004	0.0053
GR071621	1920	0.005
FR085978	1918	0.005
GR030386	1840	0.0048
ELE74009	1816	0.0048
GR056726	1784	0.0047

Number of Unique Items

```
In [186]: !cat HW3.3counts.txt | wc -l
```

12592

END STUDENT CODE HW33

HW3.3.1 OPTIONAL

Using 2 reducers: Report your findings such as number of unique products; largest basket; report the top 50 most frequently purchased items, their frequency, and their relative frequency (break ties by sorting the products alphabetical order) etc. using Hadoop Map-Reduce.

START STUDENT CODE HW331 (INSERT CELLS BELOW AS NEEDED)

END STUDENT CODE HW331

HW3.4. (Computationally prohibitive but then again Hadoop can handle this) Pairs

Suppose we want to recommend new products to the customer based on the products they have already browsed on the online website. Write a map-reduce program to find products which are frequently browsed together. Fix the support count (cooccurrence count) to $s = 100$ (i.e. product pairs need to occur together at least 100 times to be considered frequent) and find pairs of items (sometimes referred to itemsets of size 2 in association rule mining) that have a support count of 100 or more.

List the top 50 product pairs with corresponding support count (aka frequency), and relative frequency or support (number of records where they cooccur, the number of records where they cooccur/the number of baskets in the dataset) in decreasing order of support for frequent ($100 > \text{count}$) itemsets of size 2.

Use the Pairs pattern (lecture 3) to extract these frequent itemsets of size 2. Free free to use combiners if they bring value. Instrument your code with counters for count the number of times your mapper, combiner and reducers are called.

Please output records of the following form for the top 50 pairs (itemsets of size 2):

`item1, item2, support count, support`

Fix the ordering of the pairs lexicographically (left to right), and break ties in support (between pairs, if any exist) by taking the first ones in lexicographically increasing order.

Report the compute time for the Pairs job. Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers) Instrument your mapper, combiner, and reducer to count how many times each is called using Counters and report these counts.

START STUDENT CODE HW34 (INSERT CELLS BELOW AS NEEDED)

```
In [47]: %%writefile mapper3.4.py
#!/usr/bin/env python

import sys

# Set up counters to monitor/understand the number of times a map
sys.stderr.write("reporter:counter:HW3.4 Mapper Counters,Calls,1")

for line in sys.stdin:
    prods = line.split()
    # The outer loop runs through each individual item in a basket
    for item1 in prods:
        # The inner loop sorts through all of the other basket items
        for item2 in prods:
            if item1 != item2:

                # the '10' is just a placeholder because the reducer
                print '%s\t%s\t%s\t%s' % (item1, item2, 1, 10)
```

Overwriting mapper3.4.py

```

In [99]: %%writefile reducer3.4.py
          #!/usr/bin/env python

import sys
from collections import defaultdict

# Set up counters to monitor/understand the number of times a re
sys.stderr.write("reporter:counter:HW3.4 Reducer Counters,Calls,")

d = defaultdict(int)
total = 0

for line in sys.stdin:

    # the x is actually only a placeholder because the true input
    # but the output has 4
    prod1,prod2,count,x = line.split('\t',3)
    prod1 = prod1.strip()
    prod2 = prod2.strip()
    count = int(count)

    # add the count to the dictionary where appropriate
    if (prod1, prod2) in d:
        d[(prod1,prod2)] += count
    else:
        #elif (prod2, prod1) not in d:
        d[(prod1,prod2)] = count

# Calculate the total number of product interactions
for key, value in d.iteritems():
    if value >= 100:
        total += int(value)

# As mentioned in the prompt, calculate the relative frequency of
# but only for ones that have at least 100 co-occurrences
for key, value in d.iteritems():
    if value >= 100:
        freq = 0
        freq = round((float(value) / float(total)),4)
        print '%s\t%s\t%s\t%s' % (key[0], key[1], value, freq)

```

Overwriting reducer3.4.py

Hadoop Calls

In [100]: *# Hadoop commands*

```
!chmod a+x mapper3.4.py
!chmod a+x reducer3.4.py

!hdfs dfs -rm ProductPurchaseData.txt
!hdfs dfs -copyFromLocal ProductPurchaseData.txt
!hdfs dfs -rm -r HW3.4-output

!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.6.0-cdh5.8.0.jar \
  -files mapper3.4.py,reducer3.4.py \
  -reducer reducer3.4.py \
  -combiner reducer3.4.py \
  -mapper mapper3.4.py \
  -input ProductPurchaseData.txt \
  -output HW3.4-output \
  numReduceTasks 1
```

Deleted ProductPurchaseData.txt

Deleted HW3.4-output

packageJobJar: [] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.8.0.jar] /tmp/streamjob2065350177640137035.jar tmpDir=null

17/01/31 00:39:58 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032

17/01/31 00:39:58 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032

17/01/31 00:39:59 INFO mapred.FileInputFormat: Total input paths to process : 1

17/01/31 00:39:59 INFO mapreduce.JobSubmitter: number of splits :2

17/01/31 00:39:59 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1484544278544_0486

17/01/31 00:40:00 INFO impl.YarnClientImpl: Submitted application application_1484544278544_0486

17/01/31 00:40:00 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1484544278544_0486/ (http://quickstart.cloudera:8088/proxy/application_1484544278544_0486/)

3.4 ANSWER

Pairs Job Runtime: 26 seconds

My personal PC: Single PC, 4 cores, 16GB RAM, MacOS

My Virtual Machine: 3 processors, 9GB RAM, Linux

MapReduce Configuration: 2 Mappers, 1 Reducer, 2 Map Calls, 5 Reducer Calls

```
In [101]: # 3.4 OUTPUT/ANSWER
!hdfs dfs -cat HW3.4-output/part-0000* > HW3.4counts.txt
!cat HW3.4counts.txt | sort -k1,2 | head -n50
```

DAI22177	DAI62779	327	0.0017
DAI22177	DAI85309	128	0.0007
DAI22177	FR085978	106	0.0005
DAI22177	GR073461	146	0.0007
DAI22240	ELE37048	113	0.0006
DAI22896	DAI62779	261	0.0013
DAI22896	DAI75645	127	0.0006
DAI22896	GR038814	141	0.0007
DAI22896	GR073461	211	0.0011
DAI22896	SNA72163	102	0.0005
DAI23334	DAI62779	247	0.0013
DAI23334	ELE92920	140	0.0007
DAI31081	DAI62779	330	0.0017
DAI31081	DAI75645	122	0.0006
DAI31081	ELE17451	104	0.0005
DAI31081	ELE32164	265	0.0013
DAI31081	FR040251	173	0.0009
DAI31081	GR073461	159	0.0008
DAI31081	SNA80324	110	0.0006
DAI31081	DAI62779	107	0.0005

END STUDENT CODE HW34

HW3.5: Stripes

Repeat 3.4 using the stripes design pattern for finding cooccurring pairs.

Report the compute times for stripes job versus the Pairs job. Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers)

Instrument your mapper, combiner, and reducer to count how many times each is called using Counters and report these counts. Discuss the differences in these counts between the Pairs and Stripes jobs

START STUDENT CODE HW35 (INSERT CELLS BELOW AS NEEDED)

3.5 Mapper

```

In [50]: %%writefile mapper3.5.py
          #!/usr/bin/env python

          import sys
          from collections import defaultdict
          # Set up counters to monitor/understand the number of times a map
          sys.stderr.write("reporter:counter:HW3.5 Mapper Counters,Calls,1")
          d = {}
          for line in sys.stdin:
              line.strip()
              prods = line.split()

              # The outer loop runs through each individual item in a basket
              for item1 in prods:
                  if item1 not in d:
                      d[item1] = {}

                  # The inner loop sorts through all of the other basket items
                  for item2 in prods:
                      if item1 != item2:

                          # This section is creating a hash table for the pair
                          if item2 in d[item1]:
                              d[item1][item2] += 1
                          else:
                              d[item1][item2] = 1

          for key in d.iteritems():
              print "%s\t%s\t%" % (key[0], key[1], d[key[0]][key[1]])

```

Overwriting mapper3.5.py

3.5 Combiner

```

In [102]: %%writefile combiner3.5.py
          #!/usr/bin/env python

          import sys
          from collections import defaultdict
          import ast

          # Set up counters to monitor/understand the number of times a re
          sys.stderr.write("reporter:counter:HW3.5 Reducer Counters,Calls,1

          d = defaultdict(int)

          for line in sys.stdin:

              prod,stripe = line.split('\t')
              prod = prod.strip()

              # this will take the string and convert it to a dictionary
              stripe = ast.literal_eval(stripe)

              # add the count to the dictionary where appropriate
              if prod in d:
                  for key in stripe.iteritems():
                      if key[0] in d[prod]:
                          d[prod][key[0]] += stripe[key[0]]
                      else:
                          d[prod][key[0]] = stripe[key[0]]
              else:
                  d[prod] = stripe

          for key in d.iteritems():
              print '%s\t%s' % (key[0], key[1])

```

Overwriting combiner3.5.py

3.5 Reducer

```

In [103]: %%writefile reducer3.5.py
          #!/usr/bin/env python

import sys
from collections import defaultdict
import ast

# Set up counters to monitor/understand the number of times a re
sys.stderr.write("reporter:counter:HW3.5 Reducer Counters,Calls,")

d = defaultdict(int)
total = 0

for line in sys.stdin:

    prod,stripe = line.split('\t')
    prod = prod.strip()

    # this will take the string and convert it to a dictionary
    stripe = ast.literal_eval(stripe)

    # add the count to the dictionary where appropriate
    if prod in d:
        for key in stripe.iteritems():
            if key[0] in d[prod]:
                d[prod][key[0]] += stripe[key[0]]
            else:
                d[prod][key[0]] = stripe[key[0]]
    else:
        d[prod] = stripe

    # increment the total amount
    for key in d.iteritems():
        for key2, value in key[1].iteritems():
            if value >= 100:
                total += int(value)

    # calculate the frequency and print the values
    for key in d.iteritems():
        for key2, value in key[1].iteritems():
            if value >= 100:
                freq = 0
                freq = round((float(value) / float(total)), 4)
                print '%s\t%s\t%s\t%s' % (key[0], key2, value, freq)

```

Overwriting reducer3.5.py

Hadoop Calls

In [105]: *# Hadoop commands*

```
!chmod a+x mapper3.5.py
!chmod a+x reducer3.5.py
!chmod a+x combiner3.5.py

!hdfs dfs -rm ProductPurchaseData.txt
!hdfs dfs -copyFromLocal ProductPurchaseData.txt
!hdfs dfs -rm -r HW3.5-output

!hadoop jar /usr/lib/hadoop-0.20-mapreduce/contrib/streaming/hadoop-streaming-2.6.0-cdh5.8.0.jar \
  -files mapper3.5.py,combiner3.5.py,reducer3.5.py \
  -reducer reducer3.5.py \
  -combiner combiner3.5.py \
  -mapper mapper3.5.py \
  -input ProductPurchaseData.txt \
  -output HW3.5-output \
  numReduceTasks 1
```

Deleted ProductPurchaseData.txt

Deleted HW3.5-output

packageJobJar: [] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.6.0-cdh5.8.0.jar] /tmp/streamjob1532552472890765337.jar tmpDir=null

17/01/31 00:43:30 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032

17/01/31 00:43:31 INFO client.RMProxy: Connecting to Resource Manager at /0.0.0.0:8032

17/01/31 00:43:31 INFO mapred.FileInputFormat: Total input paths to process : 1

17/01/31 00:43:31 INFO mapreduce.JobSubmitter: number of splits : 2

17/01/31 00:43:31 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1484544278544_0487

17/01/31 00:43:32 INFO impl.YarnClientImpl: Submitted application application_1484544278544_0487

17/01/31 00:43:32 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1484544278544_0487/ (http://quickstart.cloudera:8088/proxy/application_1484544278544_0487/)

Pairs Job Runtime: 39 seconds

My personal PC: Single PC, 4 cores, 16GB RAM, MacOS

My Virtual Machine: 3 processors, 9GB RAM, Linux

MapReduce Configuration: 2 Mappers, 1 Reducer, 2 Map Calls, 3 Reducer Calls

```
In [106]: # 3.5 OUTPUT/ANSWER
!hdfs dfs -cat HW3.5-output/part-0000* > HW3.5counts.txt
!cat HW3.5counts.txt | sort -k1,2 | head -n50
```

DAI16732	FR078087	106	0.0002
DAI18527	SNA44451	102	0.0002
DAI22177	DAI31081	127	0.0003
DAI22177	DAI62779	382	0.0008
DAI22177	DAI63921	136	0.0003
DAI22177	DAI75645	123	0.0003
DAI22177	DAI83733	126	0.0003
DAI22177	DAI85309	172	0.0004
DAI22177	ELE17451	203	0.0004
DAI22177	ELE26917	134	0.0003
DAI22177	ELE32164	155	0.0003
DAI22177	ELE34057	107	0.0002
DAI22177	ELE56788	134	0.0003
DAI22177	ELE66600	101	0.0002
DAI22177	ELE66810	105	0.0002
DAI22177	ELE74009	108	0.0002
DAI22177	ELE91337	150	0.0003
DAI22177	FR031317	160	0.0003
DAI22177	FR032293	128	0.0003
DAI22177	FR040251	101	0.0004

END STUDENT CODE HW35

OPTIONAL

QUESTIONS BELOW THIS LINE ARE OPTIONAL

HW3.6 Computing Relative Frequencies on 100K WikiPedia pages (93Meg)

Dataset description For this assignment you will explore a set of 100,000 Wikipedia documents:

https://www.dropbox.com/s/n5lfbnztclo93ej/wikitext_100k.txt?dl=0
(https://www.dropbox.com/s/n5lfbnztclo93ej/wikitext_100k.txt?dl=0) s3://cs9223/wikitext_100k.txt, or https://s3.amazonaws.com/cs9223/wikitext_100k.txt
(https://s3.amazonaws.com/cs9223/wikitext_100k.txt) Each line in this file consists of the plain text extracted from a Wikipedia document.

Task Compute the relative frequencies of each word that occurs in the documents in wikitext_100k.txt and output the top 100 word pairs sorted by decreasing order of relative frequency.

Recall that the relative frequency (RF) of word B given word A is defined as follows:

$$f(B|A) = \text{Count}(A, B) / \text{Count}(A) = \text{Count}(A, B) / \sum_{B'} (\text{Count}(A, B'))$$

where count(A,B) is the number of times A and B co-occur within a window of two words (co-occurrence window size of two) in a document and count(A) the number of times A occurs with anything else. Intuitively, given a document collection, the relative frequency captures the proportion of time the word B appears in the same document as A. (See Section 3.3, in Data-Intensive Text Processing with MapReduce).

In the async lecture you learned different approaches to do this, and in this assignment, you will implement them:

- Write a mapreduce program which uses the Stripes approach and writes its output in a file named rfstripes.txt
- Write a mapreduce program which uses the Pairs approach and writes its output in a file named rfpairs.txt
- Compare the performance of the two approaches and output the relative performance to a file named rfcomp.txt. Compute the relative performance as follows: (running time for Pairs/ running time for Stripes). Also include an analysis comparing the communication costs for the two approaches. Instrument your mapper and reduces for counters where necessary to aid with your analysis.

NOTE: please limit your analysis to the top 100 word pairs sorted by decreasing order of relative frequency for each word (tokens with all alphabetical letters).

Please include markdown cell named rf.txt that describes the following:

HW3.7 Apriori Algorithm

What is the Apriori algorithm? Describe an example use in your domain of expertise and what kind of . Define confidence and lift.

NOTE: For the remaining homework use the online browsing behavior dataset located at (same dataset as used above):

<https://www.dropbox.com/s/zlfyiwa70poqg74/ProductPurchaseData.txt?dl=0>

Each line in this dataset represents a browsing session of a customer. On each line, each string of 8 characters represents the id of an item browsed during that session. The items are separated by spaces.

Here are the first few lines of the ProductPurchaseData FRO11987 ELE17451
ELE89019 SNA90258 GRO99222 GRO99222 GRO12298 FRO12685 ELE91550
SNA11465 ELE26917 ELE52966 FRO90334 SNA30755 ELE17451 FRO84225
SNA80192 ELE17451 GRO73461 DAI22896 SNA99873 FRO86643 ELE17451
ELE37798 FRO86643 GRO56989 ELE23393 SNA11465 ELE17451 SNA69641
FRO86643 FRO78087 SNA11465 GRO39357 ELE28573 ELE11375 DAI54444

HW3.8. Shopping Cart Analysis

Product Recommendations: The action or practice of selling additional products or services to existing customers is called cross-selling. Giving product recommendation is one of the examples of cross-selling that are frequently used by online retailers. One simple method to give product recommendations is to recommend products that are frequently browsed together by the customers.

Suppose we want to recommend new products to the customer based on the products they have already browsed on the online website. Write a program using the A-priori algorithm to find products which are frequently browsed together. Fix the support to $s = 100$ (i.e. product sets need to occur together at least 100 times to be considered frequent) and find itemsets of size 2 and 3.

Then extract association rules from these frequent items.

A rule is of the form:

$(\text{item1}, \text{item5}) \Rightarrow \text{item2}.$

List the top 10 discovered rules in decreasing order of confidence in the following format

$(\text{item1}, \text{item5}) \Rightarrow \text{item2}, \text{supportCount}, \text{support}, \text{confidence}$

HW3.8.1

Benchmark your results using the pyFIM implementation of the Apriori algorithm (Apriori - Association Rule Induction / Frequent Item Set Mining implemented by Christian Borgelt). You can download pyFIM from here:

<http://www.borgelt.net/pyfim.html> (<http://www.borgelt.net/pyfim.html>)

Comment on the results from both implementations (your Hadoop MapReduce of apriori versus pyFIM) in terms of results and execution times.

END OF HOMEWORK

